SIPHER STREET
INVESTMENT FUND

Equity Research
20 March 2026

Broadcom Inc

Industry Overview

Broadcom operates at the centre of the AI infrastructure buildout, the largest capital expenditure cycle in technology history. The five largest U.S. cloud providers have collectively guided to approximately $600 to $700 billion in capital expenditure for 2026, roughly double 2025 levels, with an estimated 75% directed toward AI infrastructure. Custom accelerators are taking share from general purpose GPUs as hyperscalers seek lower total cost of ownership for inference workloads, creating a custom ASIC market projected to reach $150 to $170 billion by 2028. Networking represents a second critical bottleneck, with spend per dollar of compute rising as cluster sizes scale into gigawatt territory. The companies best positioned in this cycle are those that can supply custom silicon, networking, and infrastructure software simultaneously, a combination that structurally narrows the competitive field.

Business Overview / Revenue Engine

Broadcom operates through two segments: Semiconductor Solutions, approximately 65% of FY25 revenue, and Infrastructure Software at 35%. Within semiconductors, the AI franchise comprises custom accelerators co developed with hyperscaler customers and networking silicon including the Tomahawk and Jericho Ethernet switching platforms that form the interconnect fabric of large-scale AI clusters. The non-AI portfolio spans wireless components anchored by Apple, broadband, enterprise networking, and server storage. Infrastructure Software is centred on VMware Cloud Foundation, the virtualisation platform positioned as the essential abstraction layer for private and hybrid cloud, supplemented by mainframe, cybersecurity, and fibre channel SAN solutions. Over 90% of software bookings are recurring, with gross margins above 93%.

We frame the revenue engine through two compounding mechanisms: customer deepening, where each hyperscaler engagement expands from a single XPU programme into multiple workload specific accelerators, networking attach, and rack scale integration over successive generations; and platform convergence, where the combination of custom silicon, networking, and VMware creates switching costs that no pure play competitor can replicate. Together they produce a trajectory where growth is driven not only by new customer wins but by expanding revenue per customer per generation.

2026 Q2-Q4 Trajectory

Q1 delivered a clean beat with revenue of $19.3 billion and AI semiconductor revenue of $8.4 billion, up 106% year over year, while gross margins held at 77% and management dismissed prior concerns around margin dilution from rack scale deployments. Q2 is guided to $22.0 billion, 7% above consensus, with AI revenue accelerating to $10.7 billion as networking reaches 40% of total AI revenue. The second half setup is favourable: Anthropic's multi gigawatt ramp, Meta's MTIA scaling, and early OpenAI volume shipments should drive meaningful sequential acceleration through Q3/Q4, consistent with management's line of sight to over $100 billion in AI chip revenue for FY27. The Q2 print in early June is the next decisive data point.

Equities

United States
Semiconductors

2-year rating	**Buy**
2y price target	**$451.54**
Current Price	**$319.21**
Upside/(Downside)	**41.5%**
TICKER (BBG):	AVGO US

Trading data and key metrics

52-wk range	$138.10-414.61
Market cap	$1.50T
Shares o/s	4,736m
Free float	98.1%
AVG. Daily Volume (L3M)	27,721k
Beta	1.25
EV / EBITDA (FY2E)	14.74x
P / E (FY2E)	17.61x

Estimates

FY (Oct)	2025A	2026E	2027E
EBITDA (mm)			
Q1	**7750A**	**11600E**	**18500E**
Q2	**7700A**	**11531E**	**19487E**
Q3	**8235A**	**15200E**	**25000E**
Q4	**9550A**	**19500E**	**29600E**
Year	**33235A**	**57831E**	**92587E**

Bhavya Patel
Lead Analyst
Sipher Street
bhavya@sipherstreet.com

Henish Patel
Lead Analyst
Sipher Street
henish@sipherstreet.com

Broadcom 2-Y Stock Chart

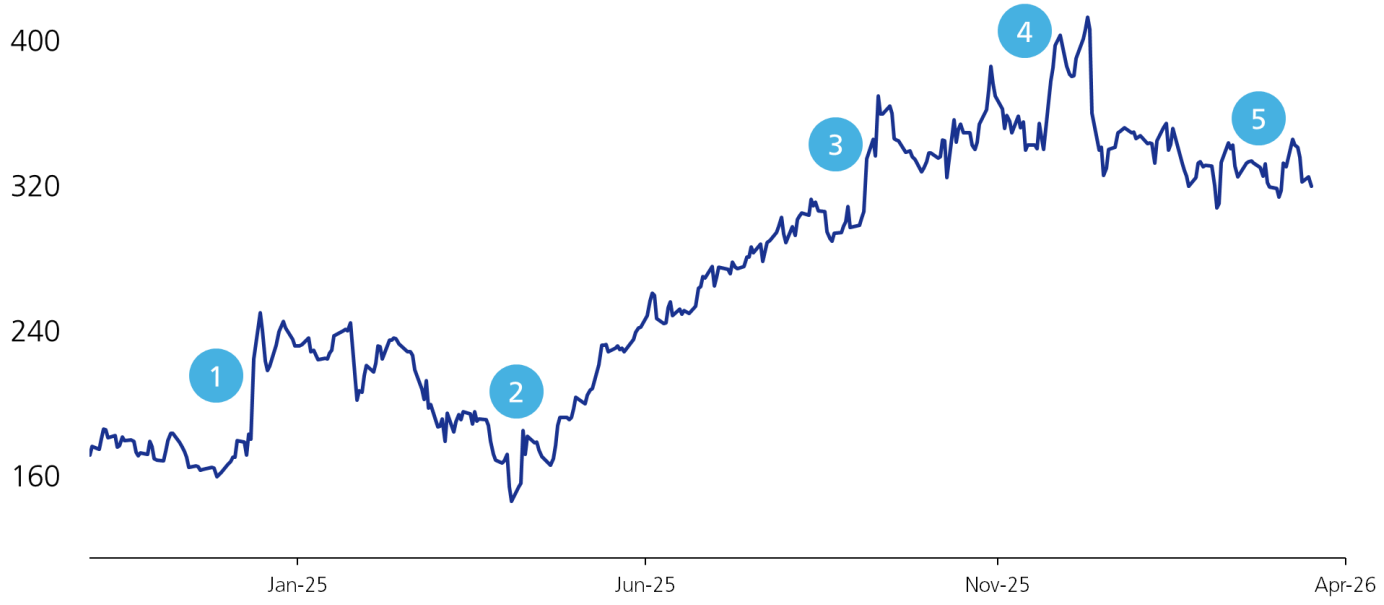

1. This move was less about the earnings beat and more about a fundamental change in how the market understood the business. Broadcom reported Q4 FY2024 on December 12 with revenue up 51% to $14.05 billion and EPS clearing consensus. But what moved the shares was CEO Hock Tan framing, for the first time, a $60-90 billion addressable market from just three hyperscaler customers by FY2027, against AI revenue of only $12.2 billion. That implied a 70-95% CAGR and AVGO surged 24% in a single session, crossing $1 trillion in market cap. Until that call, investors had treated Broadcom as a VMware integration story with an uncertain AI component. Tan replaced that uncertainty with specific commitments: million-XPU cluster deployments from all three hyperscalers and VMware software revenue nearly tripling to $21.5 billion.

2. The late-January 2025 drawdown was macro-driven, not company-specific. DeepSeek's R1 model on January 27, claiming frontier performance at a $5.6 million training cost, triggered a 17% single-session drop in AVGO and roughly $1 trillion in destruction across AI semiconductor names. The recovery came quickly: hyperscaler Q4 earnings confirmed zero capex cuts, with Google, Meta, Microsoft, and Amazon all raising 2025 spending plans, and Broadcom's own Q1 FY2025 showed AI revenue accelerating to $4.1 billion, up 77%. By April the stock had retraced the entire loss. The episode confirmed that the underlying demand trajectory was intact and that the January selloff was a sentiment event, not a structural one.

3. The August-to-October 2025 climb was driven by a compressed sequence of firm-specific wins. Q3 FY2025 on September 4 delivered record revenue of $16.0 billion with AI semiconductor revenue at $5.2 billion, up 63%, and Tan disclosed a fourth XPU customer with a $10 billion order. Weeks later, Anthropic signed a $21 billion deal for Google TPU racks, confirming it as the mystery customer, while OpenAI announced a partnership to co-develop 10 gigawatts of custom accelerators, a deal Citi estimated at $100 billion in cumulative revenue. The stock pushed toward $370 but moved in a staircase pattern, with each catalyst followed by profit-taking after a 120% run from April lows and periodic macro noise around tariffs.

4. AVGO hit its all-time high of $414.61 on December 10, the day before Q4 FY2025 earnings, with the market pricing in near-perfect execution. The results delivered: revenue of $18.0 billion up 28%, AI semiconductor revenue of $6.5 billion up 74%, and five confirmed XPU partners. Despite objectively strong numbers, the stock dropped 21% within a week. Management guided gross margins down 100 basis points due to margin-dilutive rack-level system sales, the $73 billion AI backlog came in below some expectations, and Tan's remark that 2026 demand was "hard to pinpoint" gave sellers the cover they needed after a 120% annual run. This was a de-rating from perfection, not a deterioration in the business.

5. Since January 2026, the stock has consolidated in a $307-350 range as tariff fears, extended AI chip export reviews, and a broad SaaS selloff weighed on semiconductor sentiment, with fund managers cutting chip exposure to the lowest since 2012. Q1 FY2026 on March 4 cut through that noise: revenue of $19.3 billion, AI semiconductor revenue of $8.4 billion up 106%, and Q2 guidance of $22.0 billion beating consensus by $1.6 billion. Tan confirmed a sixth XPU customer in OpenAI, disclosed line of sight to $100 billion in AI chip revenue by FY2027, and announced a $10 billion buyback. This week, Nvidia's GTC 2026 keynote reinforces the demand picture, with Jensen Huang projecting $600 billion in annual hyperscaler capex. We view the current price as a disconnect between accelerating fundamentals and depressed sentiment, and the most attractive entry point since the DeepSeek trough.

Event Type

We classify Broadcom as a quality compounder that has the added benefit of riding a secular growth inflection. The AI capex tailwind is not the thesis itself but the accelerant: the current hyperscaler buildout is compressing what would otherwise be a multi year compounding trajectory into a near term growth inflection, with management confirming secured supply chain capacity through FY28. By Q4F27, the market will have observed AI semiconductor revenue more than double year over year, validated margin stability through the XPU ramp, and de risked the customer concentration narrative as multiple hyperscalers scale to multi gigawatt deployments. Our plan is to hold through this window and begin exiting around Q4F27, by which point the compounding thesis should be substantially reflected in the share price and the marginal return on new AI capex dollars begins to diminish.

Investment Thesis

The market is pricing Broadcom as a mature semiconductor conglomerate riding a transient AI capex wave, with the stock trading approximately 25% below its December 2025 highs amid a broader rotation out of AI infrastructure names. We see the opposite. Broadcom is not a passive beneficiary of hyperscaler spending cycles; it is the indispensable design partner behind the custom silicon and Ethernet fabric that its customers have chosen over general purpose GPU alternatives, a distinction the market has not fully internalised. The combination of AI revenue doubling year over year, a $73 billion backlog, and a software franchise that provides a profitability floor no pure play semiconductor peer can match produces a business compounding at approximately 50% topline growth with a 68% adjusted EBITDA margin. At roughly 31x forward earnings on FY26 consensus EPS of approximately $10.36, the stock does not reflect the structural advantages or the duration of the growth trajectory ahead.

1. Custom AI Silicon Locks In the Highest-Value Hyperscaler Relationships for a Decade

Broadcom's custom XPU business is structurally different from transactional chip sales. Each hyperscaler engagement is a multi-year, co-design partnership where Broadcom embeds its SerDes, HBM controller, and advanced packaging IP into a chip architecture that the customer cannot easily migrate to a competitor without redesigning the entire compute stack. Google's TPU relationship, now in its seventh generation with the 3nm Ironwood chip, dates to 2016 and alone generates an estimated $8-10 billion annually. Meta's MTIA training accelerator and ByteDance's custom inference chips followed, and in the past twelve months the customer roster expanded to include Anthropic, which placed orders worth $21 billion for TPU racks, and OpenAI, which is co-developing a custom inference engine called Titan for deployment beginning in 2026. That is six confirmed hyperscaler XPU partnerships, each carrying switching costs measured in years of silicon validation and software stack integration. Marvell Technology, the closest competitor, holds an estimated 13-20% custom ASIC share against Broadcom's 70-75%, and its disclosed customer set is narrower, with Amazon's Trainium as the flagship win. In our base case, the XPU business is the primary revenue driver through FY2028E, and the depth of these co-design relationships is the clearest evidence that the market is underestimating the durability of Broadcom's AI growth.

Sell Side Analyst Estimates

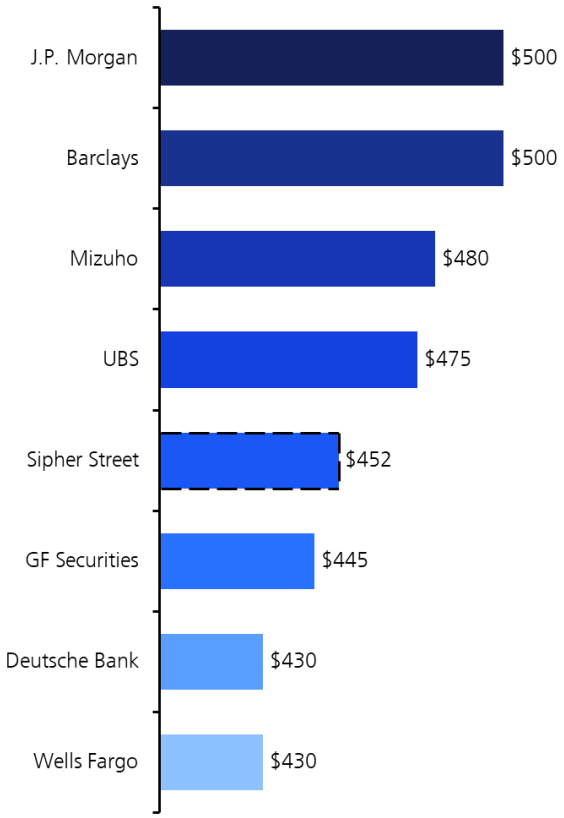

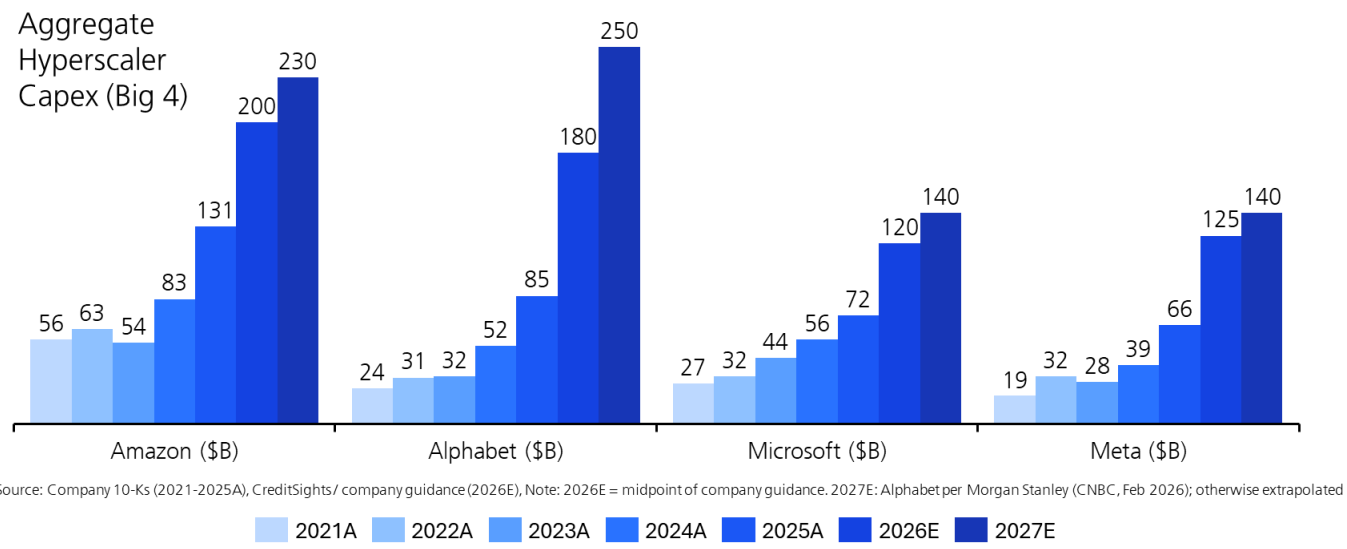

Source: Company 10-Ks (2021-2025A), CreditSights / company guidance (2026E), Note: 2026E = midpoint of company guidance. 2027E: Alphabet per Morgan Stanley (CNBC, Feb 2026); otherwise extrapolated

2. Ethernet Networking Dominance Creates a Second, Underappreciated AI Revenue Pillar

The investment narrative around Broadcom's AI business focuses disproportionately on custom accelerators, but the networking franchise is compounding at 60%+ year-over-year and now represents roughly one-third of AI semiconductor revenue, rising to an expected 40% in Q2 FY2026. Broadcom's Tomahawk 6 switch, a 102.4 Tbps single-chip solution built on 3nm process, entered volume production in early 2026, approximately a year ahead of Nvidia's competing Spectrum-X1600. The Jericho 4 router enables lossless fabric for clusters exceeding one million XPUs. The data from March 2026 confirms a critical inflection: Ethernet switch sales in AI back-end networks more than tripled in 2025, and Ethernet now commands two-thirds of AI data centre switch revenue, a dramatic reversal from InfiniBand's 80% share just two years earlier. This structural shift from Nvidia's proprietary InfiniBand protocol to open Ethernet standards is a secular tailwind that flows almost entirely to Broadcom, which holds an estimated 80–90% share of hyperscale Ethernet switching silicon. The Tomahawk 7 (204.8 Tbps) and Tomahawk 8 (409.6 Tbps) roadmap extends this leadership. In our DCF, networking revenue is the highest-confidence growth input because it is driven by cluster scale requirements that are platform-agnostic, meaning Broadcom benefits whether the customer chooses Nvidia GPUs, custom XPUs, or a hybrid architecture.

AI Back-End Network Switch Revenue Share

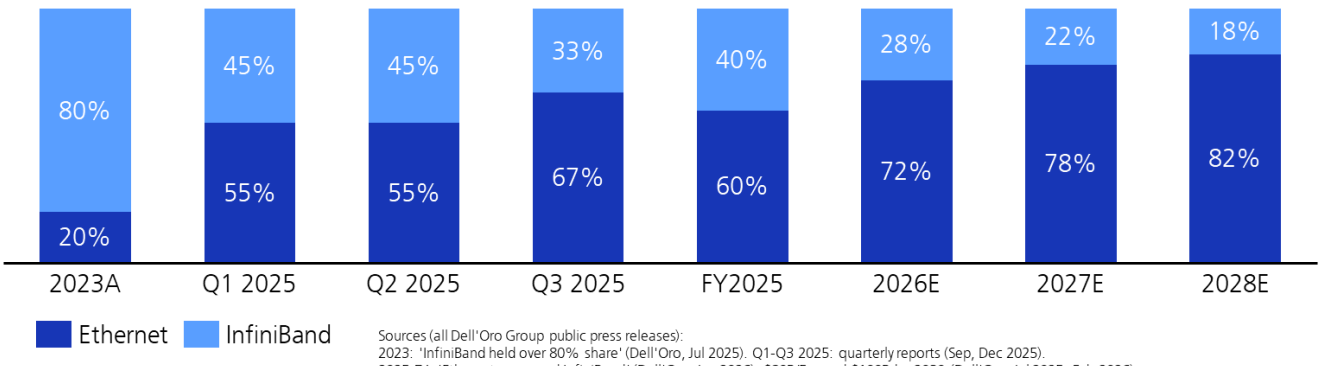

Sources (all Dell'Oro Group public press releases):
2023: 'InfiniBand held over 80% share' (Dell'Oro, Jul 2025). Q1-Q3 2025: quarterly reports (Sep, Dec 2025).
2025 FY: 'Ethernet surpassed InfiniBand' (Dell'Oro, Jan 2026). $80B/5yr and $100B by 2030 (Dell'Oro, Jul 2025, Feb 2026).
2026E-2028E: Directional estimates based on Dell'Oro commentary that Ethernet is 'projected to dominate both scale-up and scale-out' and is the 'long-term winner' (Dell'Oro, Feb 2026). Exact forward splits not publicly disclosed.

3. The VMware Software Engine Provides Profitability No Semiconductor Peer Can Match

Broadcom's infrastructure software segment, anchored by VMware and supplemented by CA Technologies and Symantec's enterprise security suite, generated $27.0 billion in FY2025 revenue at 93% gross margins and 78% operating margins, making it one of the most profitable software franchises in the technology sector. Q1 FY2026 VMware bookings reached $9.2 billion in total contract value, driving 19% annualised recurring revenue growth, and over 90% of the top 10,000 enterprise customers have now converted to subscription-based VMware Cloud Foundation licensing. The software segment's contribution to Broadcom's investment case is frequently underweighted by semiconductor-focused investors, yet it is precisely this dual-engine structure that differentiates Broadcom's risk profile from every peer. In a scenario where AI capex decelerates, the $27 billion software base still generates approximately $21 billion in annual operating income with minimal cyclicality, providing a cash flow foundation to sustain the $2.60 annual dividend, the $10 billion buyback programme, and continued debt reduction from the current $65 billion post-VMware balance. Nvidia, AMD, and Marvell have no equivalent buffer. In our base case, the software segment's margin stability is the primary reason Broadcom's downside is more contained than the market's implied volatility suggests.

Catalysts

- Hyperscaler capex revisions have only gone in one direction: The Big Five have guided $600 to $700 billion in combined 2026 capex, roughly triple the level from two years ago, and every hyperscaler reported that AI infrastructure remains supply constrained, not demand constrained. Actuals have consistently exceeded initial guidance by 30 to 45% over the past eighteen months. Each incremental capex dollar flows disproportionately to Broadcom through XPU orders, Ethernet silicon, and optical components. If the revision pattern holds, the upside to consensus estimates is material.
- The Ethernet transition is a direct revenue transfer to Broadcom: Ethernet switch sales in AI back-end networks more than tripled in 2025 and now exceed InfiniBand by a two to one ratio. Broadcom holds 80 to 90% of merchant Ethernet switching silicon, meaning the structural shift away from Nvidia's proprietary protocol to open standards flows almost entirely to one company. If current share gains continue, AI networking could rival custom XPUs as a revenue contributor by FY28.
- Next generation process nodes unlock higher ASPs across the portfolio: Broadcom has shipped the industry's first 2nm custom compute SoC on its proprietary 3.5D die stacking platform, with a family of XPU accelerators ramping in H2 2026. The 3nm Taurus 400G per lane optical DSP doubles data throughput for AI interconnects. Both carry substantially higher average selling prices than prior generations and address the most acute bottleneck in cluster scaling: bandwidth between racks.

Risks

- Customer concentration is extreme and a single defection would be material: Google alone is estimated at $8 to $10 billion annually and the top three customers likely represent over 70% of AI chip revenue. The loss of any single relationship cannot be backfilled given multi year ASIC design cycles, and the trend of hyperscalers accelerating internal silicon programmes signals that reducing third party dependency is a long term strategic priority.

- Hyperscaler capex durability beyond 2027 is the key uncertainty: The bull case requires AI infrastructure spending to compound at elevated rates through at least FY28. The four largest hyperscalers are spending 45 to 57% of revenue on capex, levels historically associated with utilities rather than technology companies, and are expected to consume approximately 90% of 2026 operating cash flow on infrastructure. If AI monetisation fails to justify these outlays or financing conditions tighten, a deceleration in 2027 to 2028 spending would directly compress Broadcom's growth trajectory.
- VMware attrition and margin mix shift could erode the software cushion: Broadcom's aggressive VMware repricing is driving measurable migration, with Nutanix reporting 2,700 former VMware customers converted in FY25 and Gartner projecting VMware could lose 35% of workloads by 2028. Separately, the shift toward rack scale AI system sales is compressing semiconductor gross margins, which dipped approximately 100 basis points to 77% in Q1F26 with further declines guided. If VMware attrition accelerates while semiconductor margins compress, the dual engine profitability thesis weakens.

Competitive Positioning & Peer Analysis

Broadcom's moat is best understood by examining how each principal competitor is structurally disadvantaged across at least one dimension that Broadcom dominates. In custom AI silicon, Marvell Technology is the only credible challenger at $8.2 billion in FY26 revenue growing 42%, but Broadcom generated $20.8 billion in AI semiconductor revenue alone in FY25, commands 70 to 75% custom ASIC share against Marvell's 13 to 20% and has six hyperscaler XPU partnerships versus Marvell's narrower set anchored primarily by Amazon. Marvell's non-GAAP gross margins of approximately 59% sit well below Broadcom's 77%, and its pipeline faces execution risk with speculation that Amazon may shift future Trainium designs to Alchip.

In AI networking, Nvidia's InfiniBand held 80% of AI back-end switch revenue as recently as 2023, but Ethernet overtook it in 2025 and now commands over two thirds of the market. This transition flows disproportionately to Broadcom, which holds 80 to 90% of merchant Ethernet switching silicon. Nvidia's Spectrum X series remains a generation behind Tomahawk 6 in volume availability. AMD competes credibly in GPU training but has no custom ASIC business, no switching silicon, and no software franchise. Intel's Gaudi accelerators have gained limited traction.

The net assessment is that no single peer replicates the combination of custom silicon, Ethernet networking dominance, and high margin VMware software. Marvell cannot match the software annuity, Nvidia cannot match the open Ethernet positioning, and AMD cannot match the custom ASIC relationships. This platform breadth is the fundamental reason Broadcom trades at a premium, and the current derating has compressed that premium to a level that does not reflect the structural advantage.

Valuation

We anchor fair value on our in-house DCF and use forward trading multiples only as a reasonableness check. The model is an unlevered FCFF framework with an explicit forecast through FY31E. Revenue is constructed from a three-pillar bottom-up build: AI semiconductors disaggregated by individual XPU customer and networking attach, non-AI semiconductors by end market, and infrastructure software split between VMware and legacy Broadcom. Gross margins expand from 70% to 75% over the forecast as the mix shifts toward higher margin networking and software, while R&D and SGA compress reflecting natural operating leverage where incremental units carry near zero marginal cost. The margin path is deliberately conservative relative to management's non-GAAP operating margins already above 66%, to build in execution risk around rack scale system deliveries and the integration of newer XPU customers. Unlevered free cash flow grows from $53.7 billion in FY26E to $113.7 billion by FY31E, with UFCF margins stabilising in the mid 50% range as the business reaches scale.

AI Semiconductor Revenue - Customer Build

USD in millions | Base Case

	2025A	2026E	2027E	2028E	2029E	2030E	2031E
Google (TPU)	14,500	30,000	48,000	56,000	60,000	62,000	63,000
% y/y		*106.9%*	*60.0%*	*16.7%*	*7.1%*	*3.3%*	*1.6%*
Anthropic	500	10,000	26,000	28,000	22,000	18,000	15,000
% y/y		*1900.0%*	*160.0%*	*7.7%*	*-(21.4%)*	*-(18.2%)*	*-(16.7%)*
Meta (MTIA)	3,000	6,000	13,000	16,000	18,000	20,000	21,000
% y/y		*100.0%*	*116.7%*	*23.1%*	*12.5%*	*11.1%*	*5.0%*
OpenAI	0	500	6,000	14,000	20,000	24,000	26,000
% y/y			*1100.0%*	*133.3%*	*42.9%*	*20.0%*	*8.3%*
ByteDance	500	1,500	3,000	5,000	6,000	6,500	6,500
% y/y		*200.0%*	*100.0%*	*66.7%*	*20.0%*	*8.3%*	*0.0%*
Apple (Baltra)	0	800	3,000	5,000	7,000	8,000	8,500
% y/y			*275.0%*	*66.7%*	*40.0%*	*14.3%*	*6.3%*
Other / Emerging	2,279	2,000	5,000	7,000	9,000	11,000	12,000
% y/y		*-(12.2%)*	*150.0%*	*40.0%*	*28.6%*	*22.2%*	*9.1%*
Total	20,779	50,800	104,000	131,000	142,000	149,500	152,000
% y/y	*70.3%*	*144.5%*	*104.7%*	*26.0%*	*8.4%*	*5.3%*	*1.7%*
AI Networking % of AI Rev	*30.0%*	*35.0%*	*37.0%*	*38.0%*	*38.0%*	*38.0%*	*38.0%*
AI Networking Revenue	6,234	17,780	38,480	49,780	53,960	56,810	57,760

The AI revenue build is the primary value driver and is constructed by individual customer programme rather than a top-down TAM approach, reflecting that Broadcom's AI business is a collection of bespoke co design partnerships, each with distinct ramp timelines and generational upgrade cycles. Google remains the anchor, scaling through the Ironwood and Sunfish TPU generations but decelerating beyond FY27E as its installed base matures and deployment growth meets power constraints. Anthropic is the highest growth customer in the near term on its 1GW to 3GW ramp, but we deliberately model it peaking in FY28E and declining thereafter to reflect the risk that a single AI lab's compute needs may plateau, a key conservatism where sell side estimates from Mizuho and Barclays are more aggressive.

OpenAI carries the widest confidence interval given its early stage first generation Titan XPU, while Meta's MTIA follows a steadier doubling trajectory supported by its disclosed $115 to $135 billion 2026 capex budget. Bytedance, Apple, and emerging customers represent optionality not fully captured in our base case. AI networking is modelled as a derived 35 to 38% of total AI revenue, anchored to management's guided range, and is our highest confidence input because it is platform agnostic: Broadcom captures networking revenue regardless of whether the customer deploys its own XPU or a competing architecture.

Non-AI Semiconductor & Infrastructure Software Revenue

USD in millions | Base Case

	2025A	2026E	2027E	2028E	2029E	2030E	2031E
NON-AI SEMICONDUCTOR							
Wireless (Apple-driven)	7,292	7,146	6,860	6,655	6,521	6,456	6,456
% y/y	*0.0%*	*-(2.0%)*	*-(4.0%)*	*-(3.0%)*	*-(2.0%)*	*-(1.0%)*	*0.0%*
Broadband	2,615	2,824	3,107	3,355	3,556	3,734	3,884
% y/y	*0.0%*	*8.0%*	*10.0%*	*8.0%*	*6.0%*	*5.0%*	*4.0%*
Server Storage Connectivity	3,771	3,884	4,039	4,161	4,286	4,371	4,459
% y/y	*0.0%*	*3.0%*	*4.0%*	*3.0%*	*3.0%*	*2.0%*	*2.0%*
Industrial & Other	625	638	657	676	690	704	718
% y/y	*0.0%*	*2.0%*	*3.0%*	*3.0%*	*2.0%*	*2.0%*	*2.0%*
Non-AI Networking (Enterprise)	2,276	2,390	2,533	2,660	2,766	2,849	2,935
% y/y	*-(40.0%)*	*5.0%*	*6.0%*	*5.0%*	*4.0%*	*3.0%*	*3.0%*
Total	16,579	16,882	17,196	17,507	17,819	18,115	18,451
% y/y	*-(8.4%)*	*1.8%*	*1.9%*	*1.8%*	*1.8%*	*1.7%*	*1.9%*
INFRASTRUCTURE SOFTWARE							
VMware	18,500	21,090	23,199	24,823	26,312	27,628	28,733
% y/y	*30.3%*	*14.0%*	*10.0%*	*7.0%*	*6.0%*	*5.0%*	*4.0%*
Legacy Broadcom Software	5,857	6,033	6,153	6,276	6,339	6,403	6,467
% y/y	*5.4%*	*3.0%*	*2.0%*	*2.0%*	*1.0%*	*1.0%*	*1.0%*
Total	24,357	27,123	29,352	31,099	32,652	34,031	35,200
% y/y	*23.3%*	*11.4%*	*8.2%*	*6.0%*	*5.0%*	*4.2%*	*3.4%*
TOTAL REVENUE	61,715	94,805	150,549	179,606	192,471	201,645	205,651
% y/y	*23.3%*	*53.6%*	*58.8%*	*19.3%*	*7.2%*	*4.8%*	*2.0%*
REVENUE MIX:							
AI Semi %	*33.7%*	*53.6%*	*69.1%*	*72.9%*	*73.8%*	*74.1%*	*73.9%*
Non-AI Semi %	*26.9%*	*17.8%*	*11.4%*	*9.7%*	*9.3%*	*9.0%*	*9.0%*
Infra Software %	*39.5%*	*28.6%*	*19.5%*	*17.3%*	*17.0%*	*16.9%*	*17.1%*

The non-AI semiconductor portfolio is modelled as a stable but structurally declining share of the mix, growing at low single digits as wireless declines modestly on Apple cycle maturity while broadband and storage benefit from enterprise refresh. We have not modelled any meaningful re acceleration, which we view as a potential source of upside. Infrastructure Software is anchored by VMware growing at 14% in FY26E decelerating to 4% by FY31E, deliberately below management's guided low double-digit growth to reflect customer attrition risk. Legacy Broadcom Software grows at 2 to 3% on a mature renewal profile. The software segment's primary role in the model is not as a growth engine but as a margin stabiliser: it provides

approximately $21 billion in annual operating income with minimal cyclicality, ensuring that even if AI capex decelerates, Broadcom's cash flow base remains sufficient to sustain its capital return programme and debt reduction.

Base Case Assumptions

The base case is built on the customer level revenue model described above, with total revenue growing from $94.8 billion in FY26E to $150.5 billion in FY27E and decelerating to $205.7 billion by FY31E at a 2.0% terminal growth rate. AI semiconductor revenue scales from $50.8 billion to $104.0 billion over FY26E to FY27E as Google, Anthropic, Meta, and OpenAI each progress through disclosed deployment milestones, before growth normalises in the outer years as gigawatt level clusters approach power and facility constraints. Gross margins expand from 70% to 75% over the forecast on mix shift toward networking and software, while R&D compresses from 15.5% to 11.5% and SGA from 5.5% to 3.8%, driving EBIT margins from 49% to approximately 60%. We apply a 20x terminal EV/EBITDA to FY31E EBITDA of $136.1 billion, producing an implied share price of approximately $452, or 41.5% upside from the current price of $319.

Bull Case Assumptions

The bull case captures an outcome where the hyperscaler capex cycle sustains at elevated rates longer than our base assumes and customer programmes ramp faster. AI semiconductor revenue reaches $67.0 billion in FY26E and $146.5 billion in FY27E, driven by Google accelerating through TPU v8 and v9 at $62.0 billion in FY27E, Anthropic scaling beyond 3GW to $36.0 billion, and OpenAI ramping its first generation XPU more aggressively at $12.0 billion. Cost of revenue compresses to 24% of sales by FY31E as higher margin networking mix and XPU IP royalties increasingly dominate, EBIT margins reach approximately 62%, and terminal year EBITDA approaches $173 billion. We apply a 22x exit multiple, a modest premium reflecting the longer duration of the growth profile in this scenario. The bull case produces an implied share price of approximately $607, or roughly 90% upside from today. This scenario requires no speculative assumptions beyond existing customer programmes ramping at the upper end of management commentary and sell side estimates from Barclays and Mizuho.

Bull Case Assumptions

The bear case models a scenario in which AI capex decelerates meaningfully beginning in FY28E as hyperscalers pull back from the current 45 to 57% capex intensity and the custom silicon TAM disappoints. AI semiconductor revenue reaches only $34.4 billion in FY26E and $69.0 billion in FY27E, roughly one third below the base case, as customer programmes scale more slowly and emerging engagements with OpenAI, Apple, and Bytedance underperform early expectations. Concurrently, VMware attrition accelerates as Nutanix and open source alternatives erode the installed base, limiting software growth to mid single digits. Gross margins settle at 72% as rack scale system sales carry a structurally heavier cost profile, R&D intensity remains elevated at 13.5% as Broadcom is forced to invest aggressively to defend its XPU position, and EBIT margins plateau at approximately 54%. We apply an 18x terminal multiple at a higher discount rate, producing an implied share price of approximately $302.91, or roughly 5.1% downside from today. That the bear case implies a contained downside even under a materially adverse AI spending scenario underscores the favourable skew: the VMware software base generating $21 billion in annual operating income provides a valuation floor that limits losses in a way that pure play semiconductor peers cannot replicate.

AVGO DCF Analysis

Implied price / Upside:	
Base case	**$451.54 / 41.5%**
Bull Case	**$607.17 / 90.2%**
Bear Case	**$302.91 / (5.1%)**

Analyst Certification

We, Bhavya Patel and Henish Patel, hereby certify that (1) all of the views expressed in this research report accurately reflect our personal views about the subject securities and issuers referred to herein and (2) no part of our compensation was, is, or will be directly or indirectly related to the specific recommendations or views expressed in this report.

Important Disclosures

Sipher Street and/or its affiliates, principals, and employees may hold positions in the securities discussed in this report and may trade in such securities from time to time without prior notice. As of the date of this publication, Sipher Street holds a long position in Broadcom Inc. (AVGO). This position may have been established before or after the publication of this report, and Sipher Street may increase, decrease, or liquidate this position at any time without obligation to notify recipients of this material.

Disclaimer

This publication has been prepared by Sipher Street for informational and educational purposes only. It does not constitute an offer, solicitation, or recommendation to purchase or sell any security or financial instrument, nor does it constitute investment advice or take into account the individual financial circumstances, investment objectives, or risk tolerance of any recipient. This material is not intended as a research report within the meaning of applicable securities regulations.

The information contained herein has been obtained from sources believed to be reliable, including company filings, publicly available data, and third party research, but Sipher Street makes no representation or warranty, express or implied, as to its accuracy, completeness, or timeliness. All estimates, projections, and forward looking statements reflect the judgment of the authors as of the date of publication, are inherently subject to significant uncertainty and risk, and are subject to change without notice. Actual results may differ materially from those expressed or implied herein. Past performance is not indicative of future results.

Sipher Street is not a registered broker dealer, investment adviser, or financial institution. Recipients of this material should conduct their own independent research and due diligence and consult with their own financial, legal, and tax advisors before making any investment decision. The value of and income from any investment may fluctuate, and investors may receive back less than the amount originally invested.

This material is provided on a confidential basis for the exclusive use of the intended recipient. No part of this publication may be reproduced, distributed, or transmitted in any form or by any means without the prior written consent of Sipher Street. Sipher Street accepts no liability for any direct, indirect, or consequential loss arising from any use of this material or its contents.

© 2026 Sipher Street. All rights reserved.